

# Temporal Pixel Masking and Conventional OCR Recovery in Short-Exposure Screen Capture: A Single-Camera UVC Study

**FIRST A. AUTHOR<sup>1</sup> [PLACEHOLDER: Author information to be filled]**

<sup>1</sup>[PLACEHOLDER: Affiliation and address to be filled]

Corresponding author: [PLACEHOLDER: Corresponding author email].

[PLACEHOLDER: Funding acknowledgment to be filled.] This manuscript includes a planned minimal-risk human-subject study. No participant data have been collected. Recruitment will begin only after a written approval, exemption, or waiver determination by the competent institutional or departmental authority; the planned safeguards are described in Section V-F.

**ABSTRACT** Screen content can be silently captured and processed by optical character recognition (OCR) or vision-language models (VLMs). This study measures conventional-OCR recovery from temporally masked screens on one USB Video Class (UVC) camera link. The pipeline partitions each frame into rapidly alternating complementary subframes and adds profile-dependent overlays. We collected 10,575 physical captures across 9 distance–angle configurations. The primary common-setting analysis uses the same 288 duplicate-averaged short-exposure units (12 content items  $\times$  8 geometries  $\times$  3 repeats) for each of three profiles. Per-capture best-of-engine mean character recovery was 94.5% unprotected, 17.8% for readability-priority, and 5.6% for high-suppression. Readability-priority retained 17.7% field-micro exact recovery across 432 annotated field opportunities, with 95.0% matched P99 character recovery. A fixed five-transform, three-engine oracle raised readability-priority recovery to 40.2% and high-suppression to 13.7%. A 150-frame mean and on-axis VLM probes recovered substantially more text. The study provides a 10,575-capture evaluation on a single UVC link and maps recovery under fixed preprocessing, temporal averaging, and VLM recognition.

**INDEX TERMS** Adversarial perturbation, camera capture, display privacy, optical character recognition, persistence of vision, screen-camera capture, temporal masking, visual eavesdropping

## I. INTRODUCTION

**S**CREENS fill offices, banks, hospitals, and conference rooms, and so do capture devices: smartphones, smart glasses [1] (e.g., Ray-Ban Meta), and miniature cameras. A single photograph of a screen suffices to extract text via optical character recognition (OCR) or identify interface elements through object detection. This *visual eavesdropping* is covert: it leaves no network traces and requires no physical contact with the target device. Controlled experiments and survey studies confirm that such attacks are common, succeed at high rates, and frequently go unnoticed [2], [3].

### A. EXISTING APPROACHES AND THEIR LIMITATIONS

Current screen-privacy solutions each address a narrow threat subset. Spatial, software, watermarking, and temporal approaches cover different viewing geometries or attack stages, but none reports the same combination of on-axis short-exposure capture, static sensitive text, and post-capture ma-


chine recognition (§II). That gap motivates measuring machine recovery from physical snapshots.

The difficulty is that degrading information for the camera also tends to degrade it for the legitimate viewer. We ask whether display temporal structure can widen this gap under short-exposure capture.

### B. CORE IDEA

Our approach exploits the gap between human temporal integration and camera exposure windows. Each frame is decomposed into  $n$  randomly complementary subframes that alternate at high speed. When the exposure is shorter than or close to one subframe duration, the camera records only an incomplete pixel subset, whereas the human observer integrates content across subframes (Fig. 1). Long exposure or video reconstruction covering a full cycle can re-sum the complementary subframes.

This study compares conventional OCR recovery across



**FIGURE 1.** Threat scenario and operating principle of temporal pixel masking. (a) A nearby camera can capture sensitive on-screen content without accessing the host system. (b) Human vision integrates complementary subframes, while a short camera exposure can intersect only part of a display cycle and leave OCR with an incomplete fragment. Long exposure or multi-frame reconstruction can accumulate more information.

three display profiles at one UVC control setting. The instrumented pipeline archives each capture as an image for offline recognition. Across the 8 common-setting geometries, matched mean recovery was 94.5% unprotected, 17.8% for readability-priority, and 5.6% for high-suppression. The high-suppression profile missed its descriptive  $< 5\%$  target.

Attack-oriented evaluation expands this result. A temporal-mean attack averaged 150 short-exposure frames, nominally 2.5 s at 60 fps, and vision-language model (VLM) probes recovered substantial large-font snippets (§V-C). These tests show where recovery rises beyond the primary conventional-OCR setting.

### C. CONTRIBUTIONS

The main contributions of this paper are:

- 1) **Controlled physical measurement:** Building on prior temporal masking ideas, we implement pixel-level complementary subframes with profile-dependent stripe/glyph overlays and a graphics processing unit (GPU) rendering prototype. The archive contains 10,575 captures from one eMeet SmartCam S600 across 3 distances and 3 angles. On 288 matched common-setting units per profile (12 content items  $\times$  8 geometries  $\times$  3 repeats), mean recovery is 94.5% unprotected, 17.8% for readability-priority, and 5.6% for high-suppression across geometrically corrected crops processed by three open-source OCR engines.
- 2) **Cluster-aware evidence and exposure sensitivity:** We report capture-level distributions and content-cluster paired contrasts. Across all 9 archived geometries, unprotected short exposure exceeds readability-priority by 78.0 percentage points (content-cluster resampling interval [75.5, 80.3]) and exceeds high-suppression by 89.1 percentage points ([85.0, 92.4]). Restricting the analysis to the 8 common-setting geometries gives cor-

responding contrasts of 76.7 points ([74.3, 79.2]) and 88.9 points ([85.5, 91.9]).

- 3) **Attack and cross-task mapping:** Per-sample strongest digital attack recovery reaches 95.2%. A 150-frame camera mean recovers 71.1% for readability-priority and 47.9% for high-suppression, rising to 79.7% and 53.9% after the underexposed geometry is excluded. VLM probes recover substantial large-font content. Detection and tracking stress tests cover four detectors on 150 real-capture COCO images and 450 MOT still-frame captures, with additional software diagnostics on 8 COCO images and 5,316 MOT frames. Together, these experiments map recovery across text recognition, temporal integration, detection, and tracking.

§II reviews prior work, §III defines the threat model, and §IV presents the design. §V reports OCR evidence, VLM probes, detection/tracking stress tests, simulation diagnostics, and the deferred usability protocol. §§VI and VII discuss the evidence hierarchy.

## II. RELATED WORK

### A. SCREEN VISUAL EAVESDROPPING THREATS

Remote screen reading attacks (telescopes, telephoto lenses) [4], deep-learning reconstruction of HDMI electro-magnetic leakage [5], covert capture by smart glasses [1], and indirect eavesdropping through reflections and refractions [6] are all documented. In a controlled Ponemon Institute study, 91% of visual-hacking attempts succeeded overall, and the first successful intrusion occurred within 15 minutes in 49% of cases [2]. Eiband et al. [3] surveyed 174 participants on the prevalence of shoulder surfing and the coping strategies users adopt. Tang and Shin [7] proposed Eye-Shield, which re-renders screen content in real time so it stays readable at the user's close viewing distance but appears blurred or pixelated at the larger distances and angles typical of shoulder surfers. Eye-Shield addresses human onlookers; our work instead measures machine recognition applied to short-exposure physical captures.

### B. DISPLAY PRIVACY PROTECTION

Physical privacy filters restrict viewing angles but cannot block an on-axis camera. Chen et al. proposed HideScreen, a display-side shoulder-surfing defense that suppresses low-frequency visual cues so screen content blends with the background beyond a designed viewing distance [8]; it targets human shoulder surfing in the spatial domain rather than short-exposure camera capture or machine recognition. Zhu et al. [9] proposed LiShield, which modulates smart LED illumination to interfere with rolling-shutter cameras, but the scheme acts on ambient lighting rather than displayed content. Visual cryptography splits a secret image into shares that combine upon overlay [10], inspiring temporal decomposition ideas, though its original formulation targets printed media. Digital watermarking [11] can embed tracking identifiers, including screen-shooting-resilient variants [12], but

provides post-hoc attribution rather than pre-emptive blocking. Software-based digital rights management (DRM) prevents digital screenshots but not physical photography.

Screen-camera visible light communication (VLC) [13]–[15] similarly exploits human temporal integration to embed information imperceptible to the eye yet decodable by cameras. Our method is VLC’s mirror image: VLC aims to let cameras *read* hidden information, whereas we aim to reduce machine readability of displayed content in constrained short-exposure capture.

Early work by Yerazunis and Carbone [16] examined temporal image masking for display privacy. Wu and Zhai [17] systematized temporal psychovisual modulation, and Gao et al. [18] applied it to information-security displays. Zhang et al. [19] proposed Kaleido, which re-encodes video through multi-frame decomposition. Rainbow [20] instead modulates ambient illumination to corrupt rolling-shutter photographs of physical objects. These schemes use imaging and subjective-quality metrics that are not interchangeable with the recovery measures used here. The cited studies do not report our specific combination of static text, short UVC exposure, and post-capture recognition. Kaleido uses chrominance-complementary frames and luminance variation, whereas our implementation uses randomized discrete pixel partitioning. The exploratory high-suppression profile also abandons strict completeness (§IV-E). Lim [21] used persistence of vision to generate anti-screen-capture keypads in a browser. Table 1 situates the present study among representative spatial, illumination, and temporal schemes.

The contribution is a system-level UVC measurement spanning explicit-field recovery, temporal integration, and VLM recognition. Detection, tracking, and digital usability proxies provide secondary diagnostics.

### C. ADVERSARIAL PERTURBATIONS AND PHYSICAL-WORLD ROBUSTNESS

Digital adversarial perturbations [22]–[24] and physical-world adversarial patches [25]–[27] are surveyed by Wei et al. [28]. Athalye et al. [29] achieved the first cross-viewpoint-robust 3D physical adversarial objects via Expectation over Transformation (EoT), though their approach still requires white-box access to the target model. Physical deployment introduces transformations absent from a digital evaluation, so robustness must be tested across capture conditions rather than inferred from pixel-space performance [25]. Physical adversarial patches exhibit cross-viewpoint robustness but require perturbation patterns pre-generated for specific classifiers [26], [27], with limited transferability to unseen models [30]. The screen-camera link also introduces moiré artifacts, which Niu et al. [31] demonstrated can themselves act as adversarial perturbations. The base mechanism partitions content temporally at the display layer rather than optimizing a model-specific spatial perturbation. Enhanced profiles also add noise and stripe/glyph overlays, whose physical transfer is evaluated empirically rather than assumed.

## III. THREAT MODEL

### A. PROTECTED ASSETS AND EVALUATED CONTENT

The motivating assets include passwords, financial data, personally identifiable information (PII), source code, and confidential documents. The corpus represents these assets through credentials, digit strings, code, paragraphs, mixed-language snippets containing embedded URL paths, and a document page. Eighteen high-value fields in 9 of 12 items were annotated before scoring; ordinary prose and incidental section numbers were excluded.

### B. ATTACKER CAPABILITIES

Attack capabilities do not form a strict monotonic hierarchy but are jointly determined by exposure, geometry, temporal sampling, and post-processing capabilities. Results are reported across the following scope:

- **Primary scope:** Instrumented snapshots at the nominal 3.91 ms UVC exposure-control setting, followed by per-frame recognition with Surya, EasyOCR, and Tesseract. The primary result uses raw camera crops. A separate three-engine stress test selects per capture among five predeclared deterministic transforms plus the raw input.
- **Extended attacks:** Long exposure, 150-frame temporal averaging, off-axis capture, phase-aware channel reconstruction, object detection and association, and commercially hosted vision-language models (§V-C).

### C. ATTACKER KNOWLEDGE

The attacker *knows* the temporal masking algorithm, the sub-frame count  $n$ , and candidate period parameters, but does *not know* the current cycle’s mask seed or phase starting point. The cryptographically secure pseudorandom number generator (CSPRNG) provides per-cycle mask randomization to prevent fixed-pattern learning and cross-cycle correlation; it does not provide a cryptographic guarantee against full-cycle averaging attacks. Once an attacker obtains and registers a complete cycle, an unknown seed cannot prevent linear reconstruction. Sparse sampling may disrupt glyph evidence used by conventional OCR, but the present design does not isolate binarization, segmentation, or recognition as the responsible stage. End-to-end VLMs recovered substantial text from the same archived fragments (§V-C), so the conventional-OCR observation does not transfer to those models.

### D. PROTECTION GOALS

We use tiered design goals under short exposure to interpret the evaluated profiles. High-suppression targets conventional OCR character recovery  $< 5\%$  with exact match  $= 0\%$ , whereas readability-priority targets a relative reduction of at least 80% with exact match  $< 1\%$ . These thresholds organize the profile comparison. Object-detection mean average precision at 0.50 intersection over union (mAP50) provides a secondary diagnostic, while  $\Delta E_{00}$  and SSIM quantify digital reconstruction quality.

**TABLE 1. Scope Comparison with Representative Display and Physical-Scene Privacy Schemes. Entries Describe Different Threats and Evaluation Targets; They Are Not Matched-Parameter Performance Claims.**

| Scheme         | Protected medium   | Target threat                                                   | Mechanism                                                                   | Reported evaluation                                                                           |
|----------------|--------------------|-----------------------------------------------------------------|-----------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------|
| HideScreen [8] | Mobile display     | Human shoulder surfing beyond a designed distance               | Spatial suppression of low-frequency cues                                   | Human-viewing distance/readability                                                            |
| LiShield [9]   | Physical scene     | Rolling-shutter camera capture                                  | Modulated ambient LED illumination                                          | Captured-image interference                                                                   |
| Kaleido [19]   | Displayed video    | Continuous-recording piracy                                     | Perceptually complete chrominance/luminance re-encoding                     | Peak signal-to-noise ratio (PSNR), structural similarity index (SSIM), and subjective quality |
| Rainbow [20]   | Physical objects   | Mobile-camera piracy                                            | Ambient-illumination modulation                                             | Rolling-shutter image degradation                                                             |
| Lim [21]       | Browser keypad     | Single digital screen capture                                   | Persistence-of-vision partial patterns                                      | Keypad capture demonstration                                                                  |
| This work      | Static screen text | Short-exposure physical capture followed by machine recognition | Randomized pixel partitioning; high-suppression profile breaks completeness | OCR/VLM recovery and integration boundaries                                                   |

## IV. SYSTEM DESIGN

### A. TEMPORAL SUBFRAME DECOMPOSITION

Given a normalized original frame  $\mathbf{I} \in [0, 1]^{H \times W \times C}$ , we generate binary masks  $\mathbf{M}_1, \dots, \mathbf{M}_n$  satisfying  $\sum_{k=1}^n \mathbf{M}_k = \mathbf{1}$ , and set  $\mathbf{I}_k = \mathbf{I} \odot \mathbf{M}_k$ . This yields:

- **Digital-domain completeness:**  $\sum_{k=1}^n \mathbf{I}_k = \mathbf{I}$ ;
- **Mutual exclusivity:** each pixel is illuminated in exactly one of the  $n$  time slots.

The distinction between “summation” and “temporal averaging” matters. On an ideal linear display with equal slot durations and equal peak luminance, the average radiance over one cycle is  $\mathbf{I}/n$ , not  $\mathbf{I}$ . Digital-domain completeness means subframes can be re-summed; it does not imply that real display luminance has been losslessly preserved.

Pixel assignments are generated from a ChaCha20 [32] CSPRNG. Rejection sampling maps the byte stream to unbiased slot indices; a separate Fisher–Yates shuffle [33] randomizes subframe playback order. A chi-square slot-count check ( $\alpha = 0.01$ ,  $\text{df} = n-1 = 3$ ) is used only as an implementation sanity check, with at most 5 resampling attempts. It is not evidence of cryptographic strength. Randomization reduces fixed patterns and cross-cycle correlation but cannot prevent linear reconstruction of a registered full cycle. Fig. 2 illustrates the complete pipeline.

### B. COMPLEMENTARY ADVERSARIAL NOISE

We overlay adversarial noise  $\mathbf{N}_1, \dots, \mathbf{N}_n$  on the subframes, satisfying:

$$\sum_{k=1}^n \mathbf{N}_k = \mathbf{0} \quad (1)$$

Equation (1) holds strictly only in the ideal linear domain prior to summation, clipping, and the display transfer function. Prior work has shown that small text-image perturbations can significantly mislead OCR systems [34], providing a basis for OCR-oriented noise design. Noise directions are generated by fast gradient sign method (FGSM)-style [22] proxy gradients with an amplitude bound of  $\varepsilon = 8/255$ ,

then decomposed complementarily in the temporal domain. Because displays cannot emit negative radiance, the implementation uses a pedestal level and numerical clipping. Consequently, after gamma correction, quantization, clipping, and camera imaging, the noise is not guaranteed to cancel physically.

**Component role:** The random dot-pattern mask is the primary temporal mechanism. Complementary noise is an optional ideal-domain component; stripe/glyph overlays and partial inversion are nonlinear additions used by the enhanced profiles below. The archived composite-profile definitions preserve correspondence with the captured data, and the physical analysis compares both individual ablations and complete profiles.


### C. LUMINANCE COMPENSATION

Each pixel is illuminated only once per  $n$ -slot cycle. A factor- $n$  light-output increase during the active slot would restore the original cycle-averaged luminance in the linear-light model. The current proof of concept (PoC) uses fixed panel output. CIEDE2000  $\Delta E_{00}$  [35] and SSIM [36] quantify digitally integrated and normalized reconstructions across configurations.

### D. TIMING AND BANDWIDTH

For a basic cycle containing only  $n$  complementary subframes, taking 60 Hz as the minimum cycle rate requires  $f_r \geq 60n$ ; at  $n = 4$  this corresponds to 240 Hz. If the cycle also includes one inversion frame, the requirement becomes  $f_r \geq 60(n+1)$ , and at 240 Hz the actual cycle rate drops to only 48 Hz. Thus, 240 Hz meets the basic four-subframe configuration but not the inversion-frame configuration’s 60 Hz cycle target. Human flicker perception also depends on luminance, contrast, retinal eccentricity, spatial content, and eye movements [37], [38]; nominal cycle rate alone cannot predict comfort.

Whether an observer integrates complementary subframes into a perceptually complete image depends on the temporal contrast sensitivity function (TCSF). Critical flicker fusion



**FIGURE 2.** Temporal pixel masking system pipeline from source content to protected display. The input frame is decomposed into CSPRNG-assigned complementary pixel masks; optional complementary OCR-oriented noise, profile-dependent stripe/glyph overlays, and an optional partial inversion frame are then inserted before high-refresh playback. A human observer is expected to integrate multiple slots into a readable view, whereas a short-exposure camera may capture only a sparse subframe. Long exposure and video temporal averaging can accumulate multiple slots and are evaluated separately.

(CFF) is not a universal cutoff: Davis et al. [38] found sensitivity near conventional rates for uniform modulation but visibility above 500 Hz when high-frequency spatial edges were present. The 48 Hz full-cycle rate of the inversion configuration may therefore produce visible flicker for some observers, and the 60 Hz basic configuration cannot be assumed flicker-free. Only the planned user study can determine whether either configuration supports comfortable temporal integration.

### E. EXPLORATORY HIGH-SUPPRESSION PROFILE

The basic linear temporal mask leaks under full-cycle temporal averaging (see §V-E2). To address this, we introduce nonlinear enhanced profiles. Unless otherwise stated, profiles in this subsection use  $n = 4$  base temporal masking with complementary adversarial noise enabled. Table 2 summarizes all profile configurations; the “Mask + noise” row in Table 6 is the base condition without stripe/glyph anti-OCR overlay.

- **Strong (anti-OCR) profile:** base mask+noise with mask cell size of 1 pixel, stripe width of 10 pixels, stripe amplitude  $\alpha_s = 0.10$ , and glyph amplitude  $\alpha_g = 0.12$ , without an inversion frame.
- **Readability-priority profile:** the Strong profile plus an  $\alpha = 0.2$  inversion frame, designated in this paper as the usability-facing composite configuration. Archived result tables and capture metadata retain the legacy label “deployed.”
- **High-suppression profile:** an exploratory security-oriented stress profile using base mask+noise with mask cell size increased to 2 pixels, stripe width of 6 pixels, stripe and glyph amplitudes of 0.42 and 0.55 respectively, plus an  $\alpha = 0.2$  inversion frame.

The nonlinear profiles depart from strict completeness ( $\sum I_k \neq I$ ), producing visible stripes and grain. This trade-

off is the opposite of Kaleido-class temporal re-encoding: the latter relies on perceptually complete mixing within a cycle to preserve viewing experience, so full-cycle integration reconstructs an approximation of the original. In the static-content capture threat addressed here, completeness is the entry point for integration attacks. The high-suppression profile therefore trades visible artifacts for lower measured recovery under full-cycle integration.

The paper-facing name “high-suppression” reflects the profile’s relative recovery in Table 2. The archived identifier remains `capture_hardened` (with legacy tag `v1m`) for reproducibility; §V-B reports the target miss and residual leakage.

### F. PARTIAL INVERSION FRAME

A long-exposure attack occurs when the exposure window covers a full display cycle ( $n = 4$  at 240 Hz yields a basic cycle of approximately 16.7 ms; with an inversion frame the  $n+1$  slot cycle is approximately 20.8 ms). The archived long-exposure labels of 31.25 and 125 ms are nominal values derived from UVC controls, not photometric measurements. In the normalized domain  $I \in [0, 1]$ , the implementation inserts a **partial inversion frame**  $\alpha(1 - I)$  after each  $n$ -subframe cycle. The equivalent 8-bit expression is  $\alpha(255 - I_8)$ . Higher  $\alpha$  changes both the integrated signal and the digital reconstruction’s luminance and contrast.

A real-capture inversion-strength ablation (Table 3; pipeline per §V-B) found 68.6% recovery at  $\alpha = 0$  and 67.0% at  $\alpha = 0.2$ . Point estimates across  $\alpha \in [0, 0.5]$  ranged from 61.7% to 72.0%; the capture-resampling intervals are descriptive and interval overlap is not an equivalence test. The ablation holds the Strong overlay fixed and uses a five-item account/code/CET6 subset. Its  $\alpha = 0.2$  value ( $N = 135$ )

**TABLE 2.** Profile Composition Summary ( $n = 4$  for All Profiles)

| Profile              | Mask | Noise | Cell (px) | Stripe width (px) | Stripe $\alpha_s$ | Glyph $\alpha_g$ | Inv. $\alpha$ |
|----------------------|------|-------|-----------|-------------------|-------------------|------------------|---------------|
| Mask only            | ✓    | –     | 1         | –                 | –                 | –                | –             |
| Mask + noise         | ✓    | ✓     | 1         | –                 | –                 | –                | –             |
| Strong               | ✓    | ✓     | 1         | 10                | 0.10              | 0.12             | –             |
| Readability-priority | ✓    | ✓     | 1         | 10                | 0.10              | 0.12             | 0.2           |
| High-suppression     | ✓    | ✓     | 2         | 6                 | 0.42              | 0.55             | 0.2           |

**TABLE 3.** Inversion Frame Strength  $\alpha$  Ablation on the Fixed Strong Overlay ( $\alpha_s = 0.10$ ,  $\alpha_g = 0.12$ ): Real-Capture Best-of-Engine Character Recovery (Mean and 95% Capture-Resampling Interval, Five Account/Code/CET6 Items Pooled Across 9 Geometric Conditions;  $N = 135$  per Setting for Long Exposure,  $N = 45$  for Video Temporal Averaging). This Five-Item Subset Differs from the 12-Item Readability-Priority Rows in Table 6.

| $\alpha$ | Long-exp. char (%) | Video avg. char (%) |
|----------|--------------------|---------------------|
| 0.0      | 68.6 [63.1, 73.9]  | 72.7 [63.3, 81.1]   |
| 0.2      | 67.0 [61.3, 72.7]  | 69.8 [59.9, 78.7]   |
| 0.3      | 72.0 [66.2, 76.9]  | 64.3 [54.0, 73.6]   |
| 0.5      | 61.7 [55.6, 67.7]  | 66.4 [56.5, 75.2]   |
| 1.0      | 18.1 [14.0, 22.0]  | 28.4 [20.5, 36.9]   |

is not a repeated estimate of the readability-priority row. Near-full inversion ( $\alpha = 1.0$ ) reduced observed recovery to 18.1% while severely degrading the digital integrated view. Under the 150-frame temporal mean, weak-inversion settings retained 64–73% recovery. The readability-priority choice of  $\alpha = 0.2$  is a modeled design decision, not a validated optimum.

Each display cycle thus contains  $n + 1$  time slots. At  $n = 4$  and 240 Hz the cycle rate is 48 Hz, below the 60 Hz design target and potentially producing perceptible flicker; no fixed “flicker-free threshold” applies universally across viewing conditions.

## G. IMPLEMENTATION SCOPE

The current PoC operates at the application layer, using Python and GPU rendering to generate display sequences. This architecture provides direct control over subframe composition, playback order, and profile parameters during evaluation.

## V. EXPERIMENTAL EVALUATION

### A. EXPERIMENTAL SETUP

**Hardware platform:** Most experiments used a desktop workstation with an Intel Core i7-8700K (3.70 GHz), 32 GB RAM, and an NVIDIA TITAN Xp (12 GB VRAM). The display was an AOC 24G51Z driven at a nominal 240 Hz. Each nominal 240-Hz slot lasts 4.17 ms, while the basic and inversion-augmented cycles are 60 and 48 Hz.

**Capture device:** All physical captures used one eMeet SmartCam S600 USB webcam with its standard UVC driver. Camera frames underwent perspective correction and content cropping before recognition. Per-capture metadata store exposure labels derived from DirectShow/UVC log<sub>2</sub> controls:

–11 and –8 map nominally to 0.49 and 3.91 ms, while –5 and –3 map to 31.25 and 125 ms. Eight geometries used the common –8/–5 settings; d0.5/15° used –11/–3. These values are control-derived labels. The display-brightness setting remained fixed, and video was captured nominally at 60 fps. The primary OCR inputs are geometrically corrected content crops.

The camera was placed at 0.5 m, 1 m, and 1.5 m from the display, rotated to 0°, 15°, and 30° at each distance. The manufacturer specifies an 8-megapixel 4K sensor and 1080p/60 fps output [39]. This grid provides nine distance–angle configurations for the physical evaluation.

**Test content:** Each position used 12 fixed items: two account strings, two code snippets, two digit strings, one English sentence, two mixed-language snippets, two English paragraphs, and one CET6 (College English Test Band 6) document. The CET6 page and two mixed-language items include Chinese characters. Three repeats give 36 captures per profile–attack–position cell and 324 across 9 positions before disclosed repeats. Detection and tracking use separate subsets as cross-task stress tests.

**Evaluation metrics:** OCR character recovery is defined as  $\max(0, 1 - \text{CER})$ , where CER is character error rate. Exact match uses case-insensitive full-text comparison after whitespace removal. Explicit-field recovery uses 18 fields fixed before scoring across 9 field-bearing content items. A field counts as recovered only when its complete Unicode-normalized form appears in at least one engine output for that capture; case, whitespace, and punctuation are ignored. We report field-level micro recovery and sample-level macro recovery separately. Leak rate is the proportion of samples with character recovery  $\geq 20\%$  and is only a descriptive threshold. On the matched common-setting units, readability-priority P50, P75, P90, P95, and P99 character recovery are 10.1%, 22.0%, 45.6%, 63.6%, and 95.0%. The corresponding high-suppression values are 3.4%, 9.2%, 15.2%, 19.1%, and 22.6%.

**OCR aggregation and uncertainty:** Character and whole-text metrics are maximized independently across the three engines for each capture. Explicit fields use the union of fields recovered by the three engines. Duplicate captures with the same profile, content, position, and repeat are averaged before matching. The primary comparison uses keys present in all three profiles. Paired contrasts resample the 12 content items as clusters with 2,000 resamples and seed 20260612. These intervals quantify sensitivity across the archived content clusters while holding geometry fixed. Table 6 gives capture-

resampling summaries for the unbalanced archive. Across all geometries, unprotected recovery exceeds readability-priority by 78.0 percentage points ([75.5, 80.3]) and high-suppression by 89.1 points ([85.0, 92.4]). Mask-only exceeds readability-priority by 3.0 points ([-0.5, 6.5]).

**Common-setting analysis:** The d0.5/15° position used a different nominal UVC setting and produced visibly underexposed protected frames. The remaining 8 geometries define the primary fixed-setting result (Table 4). After duplicate averaging and three-profile matching, each profile has  $N = 288$  units. Mean recovery is 94.5% unprotected, 17.8% readability-priority, and 5.6% high-suppression. The paired contrasts are 76.7 percentage points ([74.3, 79.2]) and 88.9 points ([85.5, 91.9]). The unbalanced captured-image mean for readability-priority is 16.7% across 408 images and is reported only as a sensitivity summary. Excluding the noncomparable geometry raises readability-priority 150-frame video-mean recovery from 71.1% to 79.7% and high-suppression from 47.9% to 53.9%.

**Additional robustness checks:** A leave-one-round-out analysis on the five repeated content items gives readability-priority short-exposure recovery of 16.6% for round 1 and 15.8% for round 2, leaving the profile-level pattern unchanged. Section V-D reports detection and tracking as cross-task stress tests.

**Fixed preprocessing stress test:** We evaluated Tesseract, EasyOCR, and Surya on all 984 selected primary-pool images. The fixed input grid comprised raw, gamma 0.5, luminance contrast-limited adaptive histogram equalization (CLAHE; 1st–99th percentile stretch, clip limit 4.0,  $8 \times 8$  tiles), unsharp masking ( $\sigma = 1.0$ , weights 1.7/–0.7), Gaussian adaptive thresholding (block 31, constant 5), and  $2 \times$  bicubic upscaling. Parameters were fixed before scoring. For each metric and capture, the attacker retained the best engine–input pair; duplicate captures were then averaged before the same three-profile matching used by the primary analysis. The complete matrix contains 17,712 cells ( $984 \text{ images} \times 6 \text{ input forms} \times 3 \text{ engines}$ ) with no OCR errors. Nonraw Tesseract cells used Tesseract 5.4.1 through `pytesseract 0.3.13`; EasyOCR 1.7.2 and Surya 0.14.7 transformed cells used CUDA on the workstation above. The resulting oracle measures a reproducible fixed-grid preprocessing attack.

Digital reconstruction proxies:  $\Delta E_{00}$  is the mean CIEDE2000 color difference between the original image and the full-cycle digitally reconstructed image; SSIM [36] is their mean structural similarity. Both metrics support within-paper comparison of digitally reconstructed profiles.

## B. REAL-CAPTURE OCR EXPERIMENT

The physical archive comprises 10,575 images (1,175 per position), including ablations, parameter sweeps, and main attack conditions. Readability-priority includes a second capture round for 5 of 12 items, giving 459/153 all-position short/video images rather than 324/108. Duplicate keys are averaged before matched contrasts. Surya [40], EasyOCR [41], and Tesseract [42] produced 31,725 engine-level

records. The Surya rerun used `surya-ocr 0.14.7`; the primary pipeline also used Tesseract 5.4.1 (`pytesseract 0.3.13`) and EasyOCR 1.7.2. Each geometry uses an archived four-point homography and output canvas, followed by the recorded content crop. The primary OCR pipeline adds no adaptive enhancement. Profiles were collected in the order original, mask-only, mask+noise, Strong, readability-priority, and high-suppression. The main comparison uses 8 common-setting geometries, while Table 6 retains all 9 as sensitivity evidence. Fig. 3 shows digital reconstructions and matched captures at 1.0 m/0°.

### Key findings:

- 1) **Primary fixed-setting evidence:** Table 4 is the main short-exposure estimate. On the same 288 matched units per profile ( $12 \text{ content items} \times 8 \text{ geometries} \times 3 \text{ repeats}$ ), mean recovery is 94.5% unprotected, 17.8% for readability-priority, and 5.6% for high-suppression. The paired content-cluster contrasts are 76.7 percentage points ([74.3, 79.2]) and 88.9 points ([85.5, 91.9]). The full 9-geometry captured-image pool gives 94.1%, 15.1%, and 5.0% across the archived settings and repeats.
- 2) **Fixed preprocessing defines a stronger OCR attack:** On the same 288 matched units per profile, the raw three-engine oracle yields 94.5% unprotected, 17.8% readability-priority, and 5.6% high-suppression recovery. Allowing the five fixed transforms raises these values to 95.9%, 40.2%, and 13.7%, respectively (Table 8). The paired unprotected-minus-readability contrast remains 55.7 percentage points ([50.3, 60.5]), while the unprotected-minus-high-suppression contrast remains 82.2 points ([78.5, 85.2]). The complete 17,712-cell matrix has no OCR errors.
- 3) **Profiles show distinct component and target behavior:** The mask-only condition reaches 19.2% in the pooled archive, while mask+noise reaches 25.6% and Strong reaches 20.7%, revealing nonmonotonic component effects. In the matched pool, readability-priority retains 17.7% field-micro exact recovery (Table 5) and 95.0% P99 character recovery. High-suppression reaches a 5.6% matched mean at the common setting, narrowly missing its descriptive  $< 5\%$  character target while meeting the exact-match = 0% target: no short-exposure capture produced an exact match (Table 6). For readability-priority, the raw common-setting reduction is  $(94.5 - 17.8)/94.5 = 81.2\%$  with 0.4% exact match, meeting the  $\geq 80\%$  relative reduction and  $< 1\%$  exact-match targets under the primary analysis. The preprocessing oracle reduces the gap to  $(95.9 - 40.2)/95.9 = 58.1\%$ , falling short of the same threshold.
- 4) **Integration attacks recover substantial text** (Fig. 4): The readability-priority profile reaches 60.9% char / 36.4% exact under nominal long exposure and 71.1% char / 42.5% exact under the 150-frame mean. Exclud-

**TABLE 4. Primary Matched Short-Exposure OCR Comparison at the Common Nominal 3.91-ms UVC Setting. Duplicate Captures Are Averaged Before Matching;  $N = 288$  per Profile Comprises 12 Content Items  $\times$  8 Geometries  $\times$  3 Repeats. The Noncomparable d0.5/15° Run Is Excluded.**

| Profile                | $N$ | Mean char recovery | Unprotected — profile | Interpretation                                             |
|------------------------|-----|--------------------|-----------------------|------------------------------------------------------------|
| Original (unprotected) | 288 | 94.5%              | —                     | Common-setting baseline                                    |
| Readability-priority   | 288 | 17.8%              | 76.7 pp [74.3, 79.2]  | Conventional-OCR reduction, with upper-tail failures       |
| High-suppression       | 288 | 5.6%               | 88.9 pp [85.5, 91.9]  | Exploratory stress profile; misses descriptive < 5% target |

**TABLE 5. Explicit Sensitive-Field Exact Recovery in the Primary Matched Short-Exposure Pool. The 18 Fields Were Fixed Before Scoring; 216 of 288 Units per Profile Contain at Least One Field.**

| Profile                | Field micro (432 opp.) | Sample macro |
|------------------------|------------------------|--------------|
| Original (unprotected) | 91.4%                  | 86.8%        |
| Readability-priority   | 17.7%                  | 17.5%        |
| High-suppression       | 1.2%                   | 1.1%         |

ing d0.5/15° raises video recovery to 79.7%. Weak inversion ( $\alpha = 0.2$ ) produces little integration reduction, and mask-only and Strong are similar under the 150-frame mean (79.0% versus 78.6%). High-suppression shows lower integration recovery, although substantial text remains recoverable.

- 5) **Long exposure produces a position-dependent reversal:** The unprotected long-exposure baseline is 47.3% char, below the readability-priority profile's 60.9%. At the sole 125 ms position (0.5 m/15°), unprotected recovery is 1.1% versus 46.3% readability-priority, with visible clipping in the full-brightness image. At 31.25 ms, the expected direction holds at 0.5 m, whereas five of six 1.0–1.5 m positions show higher readability-priority recovery. Residual stripe/glyph edges interacting with geometry-specific saturation or contrast compression provide a candidate explanation. Short-exposure recovery remains lower for readability-priority at every position.
- 6) **High-suppression reduces integration leakage:** This exploratory stress profile yields 9.3% char [7.9, 10.9] / 0% exact for nominal long exposure and 47.9% char [41.8, 54.0] / 5.6% exact [1.9, 10.2] under the 150-frame mean. Brackets are capture-resampling summaries. The pooled mean includes substantial inter-position variation and a 0% underexposed session. Excluding that session raises video recovery to 53.9% [48.0, 59.9] and exact match to 6.3%. The planned user study evaluates readability-priority rather than this visibly stronger profile.
- 7) **Geometric trends:** Across 9 conditions (0.5–1.5 m  $\times$  0–30°), protected short-exposure recovery remains below unprotected recovery. Geometry-specific values are reported descriptively because the archived design is unbalanced.

**Per-engine results:** Table 7 breaks down pooled raw-input short-exposure results by engine. The engines differ widely in unprotected recovery (Surya 37.1%, EasyOCR 79.5%,

#### Unprotected - Digital integrated reconstruction

The quick brown fox jumps over the lazy dog

#### Unprotected - Short exposure (single frame)

The quick brown fox jumps over the lazy dog

#### Unprotected - Long exposure

The quick brown fox jumps over the lazy dog

#### Readability-priority - Digital integrated reconstruction

The quick brown fox jumps over the lazy dog

#### Readability-priority - Short exposure (single frame)

The quick brown fox jumps over the lazy dog

#### Readability-priority - Long exposure

The quick brown fox jumps over the lazy dog

#### High-suppression - Digital integrated reconstruction

The quick brown fox jumps over the lazy dog

#### High-suppression - Short exposure (single frame)

The quick brown fox jumps over the lazy dog

#### High-suppression - Long exposure


The quick brown fox jumps over the lazy dog

**FIGURE 3. Real-capture montage at 1.0 m and 0°. Within each vertically stacked profile group (unprotected, readability-priority, and high-suppression), panels show, from top to bottom, the digital integrated reconstruction, 3.91 ms short-exposure capture, and 31.25 ms long-exposure capture. Camera panels are identically cropped from the archived raw camera JPEGs with no geometric correction or brightness gain, so they differ from the geometrically corrected crops used for OCR evaluation. Digital panels use the same 1920  $\times$  1080 black playback canvas and brightness-aligned full-cycle integration for each profile (for the unprotected profile this integration equals the displayed source frame); the readability-priority and high-suppression reconstructions include the recorded  $\alpha = 0.2$  inversion slot. These digital views are not direct evidence of human readability or visual comfort.**

Tesseract 84.3%). Under Strong, recovery ranges from 2.1–16.3%; under readability-priority, 3.2–11.7%; under high-suppression, 1.8–2.0%. Table 8 and Fig. 5 use the matched

**TABLE 6.** Single-UVC Real-Capture OCR Sensitivity Pool (Best-of-Engine;  $N$  = Captures per Row; All 9 Geometries). Brackets Are 95% Capture-Resampling Summaries, Not Population Confidence Intervals. Video Rows Average 150 Frames. The Primary Matched Estimate Is Table 4.

| Profile                | Attack mode            | $N$ | Char recovery [95% summary] | Exact match | Leak rate |
|------------------------|------------------------|-----|-----------------------------|-------------|-----------|
| Original (unprotected) | Short exposure         | 324 | 94.1 [93.1, 95.0]%          | 65.7%       | 99.4%     |
|                        | Video (150-frame mean) | 108 | 79.9 [73.5, 86.0]%          | 61.1%       | 86.1%     |
|                        | Long exposure          | 324 | 47.3 [42.8, 51.9]%          | 18.8%       | 58.6%     |
| Mask only              | Short exposure         | 324 | 19.2 [17.0, 21.4]%          | 0.3%        | 30.6%     |
|                        | Video (150-frame mean) | 108 | 79.0 [73.0, 84.9]%          | 48.1%       | 88.0%     |
|                        | Long exposure          | 324 | 67.2 [63.0, 71.5]%          | 35.5%       | 75.0%     |
| Mask + noise           | Short exposure         | 324 | 25.6 [22.5, 28.8]%          | 2.8%        | 36.1%     |
|                        | Video (150-frame mean) | 108 | 79.2 [73.3, 85.0]%          | 49.1%       | 88.9%     |
|                        | Long exposure          | 324 | 68.0 [63.7, 72.4]%          | 39.8%       | 76.2%     |
| Strong (anti-OCR)      | Short exposure         | 324 | 20.7 [18.2, 23.3]%          | 0.6%        | 32.4%     |
|                        | Video (150-frame mean) | 108 | 78.6 [72.5, 84.5]%          | 50.9%       | 88.9%     |
|                        | Long exposure          | 324 | 61.6 [57.1, 66.1]%          | 39.8%       | 71.3%     |
| Readability-priority   | Short exposure         | 459 | 15.1 [13.3, 17.1]%          | 0.4%        | 20.5%     |
|                        | Video (150-frame mean) | 153 | 71.1 [65.7, 76.2]%          | 42.5%       | 85.0%     |
|                        | Long exposure          | 324 | 60.9 [56.6, 65.6]%          | 36.4%       | 69.1%     |
| High-suppression       | Short exposure         | 324 | 5.0 [4.3, 5.8]%             | 0.0%        | 3.7%      |
|                        | Video (150-frame mean) | 108 | 47.9 [41.8, 54.0]%          | 5.6%        | 76.9%     |
|                        | Long exposure          | 324 | 9.3 [7.9, 10.9]%            | 0.0%        | 14.2%     |

**FIGURE 4.** Best-of-engine OCR character recovery and exact match for three representative real-capture profiles (unprotected, readability-priority, and high-suppression) under short exposure, long exposure, and a 150-frame video temporal mean, pooled across the 9 tested distance-angle configurations. Lower bars indicate less recovered text. The long-exposure and video bars expose the integration boundary: the readability-priority profile retains substantial recoverable content even when its short-exposure bar is low. Readability-priority sample counts include the disclosed five-item repeat and therefore differ from the other profiles.

common-setting pool and show that every engine benefits from at least one fixed transform. For readability-priority, recovery rises from 3.2% to 10.8% for Tesseract, 14.2% to


**TABLE 7.** Real-Capture Short-Exposure OCR: Per-Engine Results (Pooled Across 9 Geometric Conditions)

| Profile              | Engine    | Char (%) | Exact (%) | $N$ |
|----------------------|-----------|----------|-----------|-----|
| Original             | Surya     | 37.1     | 22.8      | 324 |
|                      | EasyOCR   | 79.5     | 27.8      | 324 |
|                      | Tesseract | 84.3     | 47.8      | 324 |
| Mask only            | Surya     | 4.2      | 0.3       | 324 |
|                      | EasyOCR   | 15.6     | 0.0       | 324 |
|                      | Tesseract | 3.2      | 0.0       | 324 |
| Mask + noise         | Surya     | 11.0     | 2.8       | 324 |
|                      | EasyOCR   | 18.1     | 0.0       | 324 |
|                      | Tesseract | 4.8      | 0.0       | 324 |
| Strong               | Surya     | 6.1      | 0.6       | 324 |
|                      | EasyOCR   | 16.3     | 0.0       | 324 |
|                      | Tesseract | 2.1      | 0.0       | 324 |
| Readability-priority | Surya     | 3.2      | 0.4       | 459 |
|                      | EasyOCR   | 11.7     | 0.0       | 459 |
|                      | Tesseract | 3.3      | 0.0       | 459 |
| High-suppression     | Surya     | 1.8      | 0.0       | 324 |
|                      | EasyOCR   | 2.0      | 0.0       | 324 |
|                      | Tesseract | 1.9      | 0.0       | 324 |

**TABLE 8.** Matched Common-Setting Character Recovery Under the Fixed Preprocessing Grid. Each Cell Is Raw Input  $\rightarrow$  Per-Capture Grid Oracle (%);  $N$  = 288 per Profile After Duplicate Averaging and Matching.

| Engine        | Original                | Readability             | High-supp.             |
|---------------|-------------------------|-------------------------|------------------------|
| Tesseract     | 84.9 $\rightarrow$ 91.8 | 3.2 $\rightarrow$ 10.8  | 2.0 $\rightarrow$ 8.2  |
| EasyOCR       | 85.1 $\rightarrow$ 88.4 | 14.2 $\rightarrow$ 33.9 | 2.3 $\rightarrow$ 6.5  |
| Surya         | 30.6 $\rightarrow$ 49.3 | 4.1 $\rightarrow$ 20.2  | 2.0 $\rightarrow$ 6.0  |
| Best of three | 94.5 $\rightarrow$ 95.9 | 17.8 $\rightarrow$ 40.2 | 5.6 $\rightarrow$ 13.7 |

33.9% for EasyOCR, and 4.1% to 20.2% for Surya. For high-suppression, the corresponding changes are 2.0% to 8.2%, 2.3% to 6.5%, and 2.0% to 6.0%.



**FIGURE 5. Raw-input versus fixed-grid-oracle character recovery for each OCR engine on the matched common-setting pool ( $N = 288$  per profile). The oracle selects the best of raw plus five predeclared transforms separately for each capture and metric. Lines show within-engine increases, not uncertainty intervals. Exact values appear in Table 8.**

### C. REAL-CAPTURE VLM PROBES

To test temporal masking against stronger recognizers, we applied three commercially hosted VLM endpoints exposed through the SiliconFlow application programming interface (API): Qwen3-VL-32B-Instruct [43], Kimi-K2.6 [44], and GLM-4.5V [45]. Using three models on the same photographs shows how recovery varies with model capability.

The probes cover  $0^\circ$  on-axis at three distances. The 60 fps video was split into 12 equal segments and aggregated four ways: (1) *single best*: the single frame within each segment with the highest grayscale contrast; (2) *150-frame mean*: the pixel-wise average of all 150 frames in the segment; (3) *window-mean best*: a sliding 30-frame mean, keeping the window with the highest mean luminance; and (4) *max projection*: the per-pixel temporal maximum across all frames in the segment. The 0.5 m run includes 156 inputs: 36 unprotected short-exposure baselines, 36 high-suppression short- and long-exposure inputs each, and these four aggregate views (12 segments each). The 1.0 m and 1.5 m runs each include 120 high-suppression inputs: 36 short- and long-exposure inputs each plus four video aggregate types (12 segments each). VLM inputs and the OCR best-of-engine (BoE) reference in Table 9 use the same file batch with nominal 3.91/31.25 ms labels. All VLM calls use raw camera JPEGs.

Of 1,188 model requests, 1,133 returned parseable results: 34 failures at 0.5 m, 21 at 1.0 m, and none at 1.5 m. Each table cell reports effective  $N$  and is conditional on a parseable response. GLM failures concentrate on mixed, digit, and credential items, so missingness cannot be treated as random. We computed cell-wise lower and upper bounds by assigning every failed call 0 or 100% recovery. For the GLM 0.5 m short

cell, character recovery is bounded by 61.0–74.9% and exact recovery by 44.4–58.3%. Conditional cells are not used to rank models; complete bounds for all cells are archived with the analysis.

We use the failure-free 1.5 m run as primary cross-model evidence; the 1.0 m run enters the short-exposure distance sweep as supplementary data. At 1.0 m, Qwen and Kimi long-exposure exact match are 84.4% and 74.3%, and the successful calls from the three models yield 70.0–83.3% exact match under the 150-frame video mean. GLM long-exposure exact match is only 3.3% at 1.0 m versus 58.3% at 1.5 m, but content-dependent failures preclude a cross-distance consistency claim. For compactness, Table 9 reports only the 150-frame video mean at 0.5 m; the 1.5 m run reports all four aggregate views to show view-selection sensitivity.

All calls used SiliconFlow’s OpenAI-compatible API. The recorded provider identifiers are Qwen/Qwen3-VL-32B-Instruct, Pro/moonshotai/Kimi-K2.6, and zai-org/GLM-4.5V. Qwen3-VL and GLM-4.5V have public technical reports [43], [45]. For Kimi-K2.6, the public model card [44], the provider endpoint identifier, and archived request metadata form the reproducibility record. Decoding temperature is 0 and maximum output is 4,096 tokens. The shared prompt requests only visible text and prohibits guessing; ground truth is not provided. Full prompts, parameters, and per-call outputs are archived. Metrics follow §V-A and are pooled at the capture level.

**0.5 m session selection transparency:** The archive preserves an alternative session (*real\_capture\_vlm\_d0.5\_a0\_rerun.json*, timestamp 012715) at the same position whose high-suppression short-exposure inputs are nearly all-black, yielding empty transcriptions and 0% char/exact across all three models. Table 9 uses the 141107 (3.91 ms) session because it shares the OCR reference batch and its inputs are not black. These are the only two 0.5 m on-axis VLM sessions, so Qwen’s observed short-exposure exact-match outcome ranges from 0/36 captures in the underexposed session to 28/36 captures in the tabled session, not a stable single-session estimate. The adequate-exposure criterion was applied after inspecting two sessions and is therefore a post-hoc selection rule; preserving both sessions and reporting the range makes the exposure sensitivity explicit. The tabled session is the conservative attacker-favorable case, while the alternative is retained as an underexposure diagnostic.

Table 9 presents the pooled VLM probe results: short-exposure rows show distance variation, while long-exposure and video aggregate rows show integration attack boundaries. Table 10 breaks down 0.5 m high-suppression short-exposure character recovery by content type.

#### Key findings:

- VLMs recover substantially more text than conventional OCR:** On the same 0.5 m photographs, best-of-engine OCR yields 8.4% char / 0% exact on high-suppression short exposure, while Qwen3-VL achieves 91.5% char and exactly recovers 28/36 captures. At 1.5 m, the three models retain 33.3–47.2% exact / 61.5–

**TABLE 9.** Real-Capture VLM Probes (0° On-Axis; All Rows Except the Original Short Baseline Are High-Suppression). Each VLM Cell Reports Exact Count / Char / *N* (Recoveries / % / Effective Calls); Cells with Failures Are Conditional on Parseable Responses. OCR BoE Shows the Same-Batch Three-Engine Best-of-Engine Reference (Exact / Char, %). The Temporal-Mean Rows Average 150 Short-Exposure Frames; Other Aggregate Views Are Defined in the Text. The 0.5 m High-Suppression Short Row Uses the Adequate-Exposure Attacker-Favorable 141107 Session; the Only Alternative 0.5 m Session Was Underexposed and Yielded 0/36 Qwen Exact Matches.

| Distance                             | Condition                | OCR BoE     | Qwen3-VL-32B   | Kimi-K2.6      | GLM-4.5V       |
|--------------------------------------|--------------------------|-------------|----------------|----------------|----------------|
| <i>Short-exposure distance sweep</i> |                          |             |                |                |                |
| 0.5 m                                | original short           | 58.3 / 92.4 | 30 / 94.5 / 36 | 30 / 95.6 / 36 | 19 / 92.0 / 26 |
| 0.5 m                                | short                    | 0.0 / 8.4   | 28 / 91.5 / 36 | 17 / 84.7 / 36 | 16 / 70.8 / 31 |
| 1.0 m                                | short                    | 0.0 / 3.6   | 20 / 79.1 / 36 | 14 / 68.6 / 36 | 9 / 47.0 / 33  |
| 1.5 m                                | short                    | 0.0 / 6.5   | 17 / 68.9 / 36 | 12 / 61.5 / 36 | 13 / 70.0 / 36 |
| <i>Key attack views</i>              |                          |             |                |                |                |
| 0.5 m                                | long                     | 0.0 / 2.1   | 28 / 85.7 / 34 | 1 / 7.9 / 36   | 0 / 4.9 / 30   |
| 0.5 m                                | Video (150-frame mean)   | 8.3 / 50.7  | 10 / 94.0 / 12 | 10 / 95.5 / 12 | 7 / 91.8 / 8   |
| 1.5 m                                | long                     | 0.0 / 13.2  | 33 / 89.3 / 36 | 32 / 89.9 / 36 | 21 / 62.4 / 36 |
| 1.5 m                                | Video (single best)      | 0.0 / 4.3   | 1 / 24.8 / 12  | 0 / 12.7 / 12  | 0 / 29.8 / 12  |
| 1.5 m                                | Video (150-frame mean)   | 0.0 / 48.4  | 8 / 90.3 / 12  | 8 / 90.3 / 12  | 7 / 88.5 / 12  |
| 1.5 m                                | Video (window-mean best) | 0.0 / 8.6   | 7 / 89.0 / 12  | 5 / 84.2 / 12  | 8 / 87.2 / 12  |
| 1.5 m                                | Video (max projection)   | 0.0 / 19.0  | 8 / 90.1 / 12  | 7 / 89.2 / 12  | 7 / 88.7 / 12  |

**TABLE 10.** 0.5 m High-Suppression Short-Exposure VLM Character Recovery by Content Type. Baseline Column Shows Same-Session Unprotected Original Short Results for Qwen (Kimi/GLM Similar); GLM Had Occasional Call Failures, Parenthesized *N* Shows Effective Count. At 1.5 m, CET6 Dense Pages Yield 0% Char Across All Three Models.

| Content type ( <i>N</i> ) | Baseline (%) | Qwen (%) | Kimi (%) | GLM (%)  |
|---------------------------|--------------|----------|----------|----------|
| CET6 document (3)         | 64.4         | 18.5     | 5.6      | 0.4      |
| Credentials (6)           | 97.6         | 97.6     | 88.3     | 77.1 (5) |
| Code snippets (6)         | 89.8         | 94.9     | 95.8     | 88.9     |
| Digit strings (6)         | 100.0        | 100.0    | 98.6     | 50.0 (4) |
| Mixed content (6)         | 100.0        | 100.0    | 92.5     | 98.2 (4) |
| Paragraphs (6)            | 100.0        | 99.8     | 81.7     | 66.5     |
| English sentences (3)     | 95.3         | 95.3     | 96.9     | 94.6     |

70.0% char. The gap between conventional OCR and modern VLMs is large.

- VLM recovery concentrates on large-font short snippets:** At 0.5 m, CET6 document pages yield 0.4–18.5% VLM character recovery, whereas credentials, code, and digit strings yield 50–100% (Table 10). At 1.5 m, CET6 pages drop to 0% across the three models, while digit strings and short sentences remain recoverable. The pattern is consistent with font scale and language priors [46], [47].
- Attack capability varies across models and conditions:** In the failure-free 1.5 m session, nominal long-exposure exact recovery ranges from 58.3% to 91.7%, with 62.4–89.9% character recovery. Endpoint, distance, capture view, and content jointly shape this spread.
- Sustained video aggregation strongly increases VLM recovery:** At 0.5 m, missingness bounds for the 150-frame mean span 58.3–91.7% exact and 61.2–95.5% character recovery across models. In the failure-free 1.5 m session, the 150-frame mean yields 58.3–66.7% exact and 88.5–90.3% character recovery, while single-best-frame exact recovery is 0–8.3%. The ac-

companying 12.7–29.8% single-best character values rest on low-information inputs, where fluent language-prior completions can inflate character recovery even though the prompt prohibits guessing; the exact-match counts are the more conservative reading.

High-suppression strongly reduces dense small-text recovery under conventional OCR, but VLMs still recover many large-font credentials, code snippets, and digit strings. Content-adaptive, VLM-aware temporal protection is the logical next step.

#### D. REAL-CAPTURE OBJECT DETECTION AND TRACKING

To probe whether the temporal degradation transfers beyond text, we report an exploratory detection/tracking stress test within the same camera link; evidence strength is not equivalent to the OCR main experiment. Content was displayed on the 240 Hz screen, captured by the S600 at 1.5 m / 0°, and processed by four detectors [48]–[51]. mAP, mAP50, and average recall (AR) are reported for detection. Tracking uses higher-order tracking accuracy (HOTA), identification F1 (IDF1), multiple-object tracking accuracy (MOTA), and multiple-object tracking precision (MOTP). The protected condition uses the  $n = 4$  mask+noise profile without overlays or inversion. The `real_clean` reference is an unprotected nominal 31.25-ms capture. The `real_short` image is the maximum-grayscale-contrast frame from a three-frame protected burst at the nominal 3.91-ms setting, and `real_video` is the mean of 150 protected frames acquired nominally at 60 fps. Tracking uses 450 frames from MOT17-09-FRCNN via still-display-then-capture and an in-project greedy association implementation inspired by ByteTrack [52]. Protection, exposure, resolution, and frame reduction differ across conditions.

Under short exposure, all four detectors' mAP50 drops from 39.2–50.2% to 1.6–5.8% (Fig. 6), relative decreases exceeding 85%. The `real_video` condition yields 3.0–7.3%,

**TABLE 11.** Real-Capture COCO Detection (150 Images; *real\_clean* = Unprotected Capture Baseline)

| Detector     | Condition         | mAP   | mAP50 | AR    |
|--------------|-------------------|-------|-------|-------|
| YOLO26x      | <i>real_clean</i> | 28.6% | 43.9% | 31.0% |
|              | <i>real_short</i> | 1.3%  | 2.4%  | 1.4%  |
|              | <i>real_video</i> | 3.5%  | 5.3%  | 3.6%  |
| RT-DETR-x    | <i>real_clean</i> | 30.1% | 50.2% | 34.8% |
|              | <i>real_short</i> | 3.5%  | 5.8%  | 4.1%  |
|              | <i>real_video</i> | 4.6%  | 7.3%  | 5.5%  |
| Faster R-CNN | <i>real_clean</i> | 21.3% | 39.2% | 28.7% |
|              | <i>real_short</i> | 0.8%  | 1.6%  | 0.9%  |
|              | <i>real_video</i> | 1.5%  | 3.0%  | 1.8%  |
| RetinaNet    | <i>real_clean</i> | 22.6% | 39.4% | 27.0% |
|              | <i>real_short</i> | 0.9%  | 1.8%  | 1.0%  |
|              | <i>real_video</i> | 1.9%  | 3.1%  | 2.0%  |

**TABLE 12.** Real-Capture MOT17-09-FRCNN Still-Frame Tracking (ByteTrack-Style Greedy Association, 450 Frames)

| Detector     | Condition         | HOTA  | IDF1  |
|--------------|-------------------|-------|-------|
| YOLO26x      | <i>real_clean</i> | 15.8% | 12.1% |
|              | <i>real_short</i> | 8.7%  | 6.5%  |
|              | <i>real_video</i> | 10.0% | 8.7%  |
| RT-DETR-x    | <i>real_clean</i> | 14.2% | 10.4% |
|              | <i>real_short</i> | 8.1%  | 5.2%  |
|              | <i>real_video</i> | 9.6%  | 6.8%  |
| Faster R-CNN | <i>real_clean</i> | 14.4% | 9.6%  |
|              | <i>real_short</i> | 6.8%  | 6.0%  |
|              | <i>real_video</i> | 10.5% | 7.1%  |
| RetinaNet    | <i>real_clean</i> | 14.2% | 9.8%  |
|              | <i>real_short</i> | 5.1%  | 4.4%  |
|              | <i>real_video</i> | 6.8%  | 5.9%  |

also below the *real\_clean* baselines. Because protection, exposure, resolution, and frame reduction differ across conditions, these values form a cross-task stress test rather than an isolated protection effect. In still-frame tracking, HOTA decreases from 14.2–15.8% to 5.1–8.7% and IDF1 from 9.6–12.1% to 4.4–6.5%.


### E. SOFTWARE SIMULATION EXPERIMENTS

This section uses software simulation to address two questions: (1) what is the machine readability of one protected subframe obtained after the application renderer? (2) what is recovery under zero camera noise and perfect alignment? The experiment characterizes the post-render single-subframe capture point. Malware that intercepts the original frame or accumulates a complete cycle bypasses this path.


#### 1) Synthetic Single-Subframe OCR

Using Tesseract [42], EasyOCR [41], and Surya [40] on 120 synthetic samples, character recovery drops from 94.0–97.0% to 0–2.6% across three engines (Table 13). Stratified breakdowns further characterize recovery across the evaluated content categories.

Fig. 7 visualizes these results. Single-subframe recovery does not exceed 2.6% for any engine, confirming a consistent

**FIGURE 6.** Real-capture mAP50 for four detectors across the unprotected capture, short-exposure protected capture, and 150-frame protected mean. Conditions differ in protection, exposure, resolution, and frame reduction.**TABLE 13.** Synthetic Single-Subframe OCR: Three-Engine Character Recovery (120 Samples,  $n = 4$ ,  $\epsilon = 8/255$ )

| OCR engine | Original frame | Single subframe | Reduction |
|------------|----------------|-----------------|-----------|
| Tesseract  | 94.0%          | 0.0%            | 94.0%     |
| EasyOCR    | 94.1%          | 0.0%            | 94.1%     |
| Surya      | 97.0%          | 2.6%            | 94.4%     |

**FIGURE 7.** Character recovery for three OCR engines on original frames and single protected subframes; error bars are 95% bootstrap intervals resampled by corpus sample.

pattern across all three recognizers.

#### 2) Theoretical Security Boundary

When an attacker accumulates registered frames covering a full cycle, complementary subframes can be re-summed in the ideal linear model of (1). The archived per-sample strongest digital attack recovers 95.2%, while pure full-cycle temporal averaging recovers 94.3%. These values empirically

**TABLE 14. Archived Digital Reconstruction Boundary (Tesseract)**

| Setting                    | $N$ | Char. | Exact | SSIM  | $\Delta E_{00}$ |
|----------------------------|-----|-------|-------|-------|-----------------|
| Base: cycle mean           | 120 | 94.3% | 85.8% | —     | —               |
| Base: per-sample strongest | 120 | 95.2% | 91.7% | —     | —               |
| Readability: cycle mean    | 16  | 87.9% | 50.0% | 0.948 | 1.46            |
| High: cycle mean           | 16  | 56.6% | 6.3%  | 0.764 | 4.87            |
| High: inversion slot       | 16  | 93.0% | 81.2% | —     | —               |

The 120-sample base attack and 16-sample profile ablation are separate archived digital evaluations. “Per-sample strongest” selects the best result among seven prespecified reconstructions (temporal average cycle, blue channel max, differential blue, differential luma, global shutter slot 0, phase search max, and phase search mean) for each sample. “Inversion slot” extracts only the  $\alpha(1 - I)$  frame and applies contrast stretching; its 93.0% recovery identifies slot-level composition as a critical design variable. Profile SSIM and  $\Delta E_{00}$  describe full-cycle reconstruction quality, not physical or user-perceived quality.

**TABLE 15. COCO val2017 Simulated Detection Pipeline Diagnostics (4 Detectors, 8 Images)**

| Detector     | Condition       | mAP   | mAP50 | AR    |
|--------------|-----------------|-------|-------|-------|
| YOLO26x      | clean           | 48.1% | 58.3% | 51.0% |
|              | single subframe | 0.0%  | 0.0%  | 0.0%  |
|              | temporal avg.   | 14.3% | 18.3% | 14.3% |
| RT-DETR-x    | clean           | 47.1% | 58.7% | 53.4% |
|              | single subframe | 5.2%  | 6.0%  | 5.2%  |
|              | temporal avg.   | 16.8% | 23.7% | 17.9% |
| Faster R-CNN | clean           | 38.2% | 52.7% | 43.6% |
|              | single subframe | 3.0%  | 4.4%  | 3.0%  |
|              | temporal avg.   | 11.1% | 16.5% | 11.1% |
| RetinaNet    | clean           | 33.2% | 47.1% | 35.3% |
|              | single subframe | 3.3%  | 4.2%  | 3.3%  |
|              | temporal avg.   | 9.6%  | 13.2% | 9.6%  |

mark the reconstruction boundary of the evaluated linear decomposition. The high-suppression profile yields 47.9% pooled recovery under the 150-frame mean and 53.9% after excluding the underexposed position. Character recovery also depends on content and the physical imaging path.


### 3) Detection and Tracking Simulation

Simulations were run on COCO val2017 [53] and MOT17 [54] with YOLO26x [48], RT-DETR-x [49], Faster R-CNN [50], and RetinaNet [51]. The COCO simulation used 8 images as an implementation pipeline diagnostic (Table 15). The MOT simulation covers 7 sequences totaling 5,316 frames and uses the same in-project ByteTrack-inspired greedy association method (Table 16).

On the 8 COCO images, single-subframe mAP50 across the four detectors is 0–6.0%. Across the 5,316 MOT frames, single-subframe IDF1 is 0.1–5.4%, while MOTA includes negative and nonmonotonic values. Negative MOTA occurs when the false-positive, false-negative, and identity-error penalties exceed the number of ground-truth objects; it is not a negative detection probability. Detector architecture, confidence thresholds, proposal density, and the greedy associator jointly shape the observed differences.

**TABLE 16. MOT17 Simulated Tracking (ByteTrack-Style Greedy Association; 5,316 Frames)**

| Detector     | Condition       | MOTA   | MOTP  | IDF1  |
|--------------|-----------------|--------|-------|-------|
| YOLO26x      | clean           | 31.1%  | 81.7% | 27.1% |
|              | single subframe | 0.1%   | 74.8% | 0.1%  |
|              | temporal avg.   | 9.1%   | 78.1% | 12.0% |
| RT-DETR-x    | clean           | 24.1%  | 79.8% | 38.2% |
|              | single subframe | −2.3%  | 73.2% | 5.4%  |
|              | temporal avg.   | −11.4% | 75.0% | 14.7% |
| Faster R-CNN | clean           | 3.8%   | 77.9% | 35.6% |
|              | single subframe | −2.5%  | 71.2% | 4.5%  |
|              | temporal avg.   | −2.7%  | 73.2% | 14.0% |
| RetinaNet    | clean           | −5.4%  | 77.1% | 30.5% |
|              | single subframe | 0.4%   | 70.6% | 2.7%  |
|              | temporal avg.   | −2.9%  | 73.8% | 11.9% |

**FIGURE 8. Residual capability of one protected subframe relative to the clean baseline (clean = 100%) across the evaluated OCR, detection, and tracking pipelines. The metrics are normalized within task and are not directly commensurate across tasks.**

### F. USER EXPERIENCE STUDY

To evaluate the impact of the protected display on human readability and visual comfort, we designed a web-based user study.

**Equipment and environment protocol:** The same 240 Hz display as in §V-A. The formal study version sets a hard measured refresh-rate threshold of  $\geq 200$  Hz to ensure the fixed  $n = 4$  basic subframe cycle rate  $f_r/n \geq 50$  Hz; below this threshold participation is blocked rather than silently reducing  $n$ . The 144–199 Hz adaptive path is available only for an explicitly labeled demonstration mode, and its sessions are excluded from analysis by default. Before each participant begins, the experimenter confirms fixed display brightness, disabled auto-brightness/power-saving and variable refresh rate, unified browser fullscreen, closed high-load background processes, and an approximate 60 cm viewing distance. Refresh rate is re-measured after any display-mode change.

The browser records actual frame intervals in each trial’s

requestAnimationFrame loop; if the interval between adjacent callbacks exceeds  $1.5\times$  the expected interval computed from the pre-trial measured refresh rate, a dropped-frame event is logged, along with estimated drop count, drop rate, mean/max frame intervals, observed refresh rate, basic subframe cycle rate, and full cycle rate with inversion frame. This recording reflects only browser presentation cadence and cannot substitute for panel scanning or photometric measurement.

**Ethics and safety:** The planned study involves minimal-risk human-subjects research, and no participant data have been collected. Before recruitment, the authors will obtain and document a written approval, exemption, or waiver determination from the competent institutional or departmental authority; formal data collection will not begin without that determination. Participants must separately check informed consent and photosensitivity screening: the page explicitly states that the masked display involves rapid temporal modulation, that participants should stop immediately upon experiencing discomfort, eye fatigue, dizziness, nausea, or headache, and that they may withdraw unconditionally at any time. Individuals who self-report a history of photosensitive epilepsy or sensitivity to flicker are not permitted to proceed. The system records consent timestamps and screening pass flags. The implementation applies a conservative  $f_r/n \geq 50$  Hz operating floor, but this engineering gate is not an ethical or clinical safety determination. At  $n = 4$  on a 240 Hz panel the basic cycle rate is 60 Hz; the inversion frame reduces the full cycle frequency to  $f_r/(n + 1) = 48$  Hz, and the implementation separately logs warnings and actual timing. Registration information (student ID, name) is used only for participation management; analysis and publication use only de-identified data. Age and gender are optional demographic fields.

**Task design:** Each participant first completes one 12-second, unscored unmasked Canvas typing warm-up, views a 10-second, unscored readability-priority full-mask preview, and then completes an 8-second, unscored typing practice under the same readability-priority mask. This sequence separates initial stimulus and input adaptation from the four subsequent 20-second scored trials (within-subjects design). On every typing screen, including warm-up, practice, and scored trials, the stimulus Canvas is visible and its player has started before the participant clicks “Start.” Clicking “Start” triggers a 5-second countdown while the input remains disabled; the timed typing interval and input availability begin only when the countdown ends. The same pre-start viewing procedure applies to both control and masked conditions.

The two conditions are each repeated twice. The experimenter’s continuously assigned zero-based sequence number  $k$  selects ABBA or BAAB order to counterbalance practice and fatigue: (A) **original text** (control), displayed as unmasked Canvas with  $n = 1$ ; (B) **masked text**, using the readability-priority configuration:  $n = 4$ , mask+noise, strong anti-OCR ( $\alpha_s = 0.10$ ,  $\alpha_g = 0.12$ , 6 cycles) + weak inversion ( $\alpha = 0.2$ ). The typing order index is  $k \bmod 2$ , and the

subjective-rating Latin-square row is  $[k/2] \bmod 6$ . Every consecutive 12 participants therefore cover all  $2 \times 6$  joint assignments. Formal sessions enforce a uniqueness constraint on  $k$ , and the backend recomputes and verifies both order indices.

Both conditions are rendered by the same `renderSourceCanvas` path, with identical dark background, light text, canvas size, 23 px bold font, and color. The sole processing difference is the temporal protection pipeline, avoiding polarity and font-weight confounds. Here “cycles” means that the web implementation rotates through six distinct mask, noise, and stripe instances before repeating. This web-playback-specific setting shares  $\alpha_s$  and  $\alpha_g$  with the real-capture configuration in §IV and avoids repeatedly integrating one fixed overlay pattern.

Scoring no longer compares characters position-by-position but aligns the full transcription against the optimal target prefix using Levenshtein minimum string distance (MSD) [55]; untyped target suffix does not count as error. The alignment yields accuracy, characters per minute (CPM), and words per minute (WPM), along with edit distance, aligned prefix length, and scoring version. First-keystroke latency is defined from the end of the 5-second countdown, when the input is enabled and the scored timer starts, to the first input event that makes the transcription non-empty. Analysis first averages each participant’s two trials per condition, then performs within-subject paired comparison.

**Subjective ratings:** After the typing trials, participants rate 6 conditions on the three 1–5 Likert items displayed by the interface: readability (1 = difficult to read, 5 = very clear), stability (stored in the `flicker` field; 1 = strong flicker, 5 = no perceptible flicker), and immediate visual comfort (`fatigue`; 1 = very uncomfortable, 5 = very comfortable). Higher values therefore indicate a better outcome on every item. Conditions include the unmasked original anchor,  $n \in \{2, 3, 4\}$  mask+noise,  $n = 4$  mask-only, and the same  $n = 4$  readability-priority full configuration (mask+noise+strong anti-OCR+weak inversion) used in the typing task. The first three masking levels all achieve  $\geq 60$  Hz real temporal playback on a 240 Hz panel; the originally planned  $n = 8$  would yield 30 Hz and trigger static fallback, so it is not included as a temporal ablation. The six conditions are presented in a 6th-order balanced Latin square; each condition must be viewed for at least 10 seconds before submission is enabled, with viewing start/end times and duration recorded. The unmasked anchor is also an attention/scale-interpretation aid but is not mechanically excluded by a “must rate 5” rule. This study does not include the high-suppression profile; assessment of its visible stripes and grain’s acceptability is left to future work (see §VI).

**Sample size and analysis plan:** A minimum of  $N = 24$  participants is planned, a number within the range required for paired effect sizes of  $d_z \approx 0.6$ – $0.8$  at 80% power, and exactly covering two rounds of the 12-person joint order assignment. Pre-registered exclusion criteria are debug/demo sessions, incomplete four typing or six rating records, mea-

**TABLE 17. Typing Task: Within-Subject Comparison of the Control and Readability-Priority Masked Conditions ( $N = XX$ ; Participant Means Over Two Trials per Condition; ' = Masked – Control With Participant-Level Bootstrap 95% CI). Per Metric, a Paired  $t$ -Test Is Used When the Pre-Specified Shapiro Test Passes (Effect Size Cohen's  $d_z$ ), Otherwise a Wilcoxon Signed-Rank Test (Effect Size Rank-Biserial  $r_{tb}$ ); Test Statistics Are Reported in the Text**

| Metric                 | Control | Masked | $\Delta$ [95% CI] |
|------------------------|---------|--------|-------------------|
| Accuracy (%)           | –       | –      | – [–, –]          |
| WPM                    | –       | –      | – [–, –]          |
| CPM                    | –       | –      | – [–, –]          |
| Attempted chars        | –       | –      | – [–, –]          |
| First-key latency (ms) | –       | –      | – [–, –]          |


**TABLE 18. Subjective Ratings for the Six Display Conditions ( $N = XX$ ; 1–5 Likert, Higher Is Better; Mean  $\pm$  SD). Stability Is the Inverse-Flicker Item and Comfort the Immediate Visual Comfort Item (§V-F); Friedman and Holm-Corrected Pairwise Wilcoxon Statistics Are Reported in the Text**

| Condition                           | Readability | Stability | Comfort |
|-------------------------------------|-------------|-----------|---------|
| Unmasked anchor                     | –           | –         | –       |
| $n=2$ mask+noise                    | –           | –         | –       |
| $n=3$ mask+noise                    | –           | –         | –       |
| $n=4$ mask+noise                    | –           | –         | –       |
| $n=4$ mask-only                     | –           | –         | –       |
| Readability-priority full ( $n=4$ ) | –           | –         | –       |


sured refresh rate below 200 Hz, any typing trial with fewer than 5 attempted characters, control mean accuracy below 50%, masked trials not in temporal mode, any trial with observed basic cycle rate below 50 Hz, or minimum-view straight-line ratings. The last criterion is defined as all six rating rows having a recorded viewing duration no greater than 11 seconds (a 1-second tolerance above the 10-second submission gate) and, within every row, identical scores on all three items. WPM, CPM, accuracy, attempted characters, and first-keystroke latency are first averaged within each participant by condition; speed metrics must be interpreted jointly with accuracy to avoid misinterpreting fast input under low accuracy as performance improvement. Paired differences are tested using paired  $t$ -tests when the pre-specified Shapiro test is passed, otherwise Wilcoxon signed-rank tests, with effect sizes and participant-level bootstrap 95% confidence intervals reported. Likert data are analyzed per dimension using Friedman tests with Holm-corrected post-hoc pairwise Wilcoxon tests. Analysis scripts generate de-identified participant means, JSON audit reports, and two L<sup>A</sup>T<sub>E</sub>X tables directly from the formal database `study_formal.db`.

**Participants:** XX volunteers were recruited; XX sessions met all pre-registered inclusion criteria and entered analysis (XX excluded; per-criterion audit counts accompany the released analysis output). Optional demographics for the included sample: XX of XX reported age (median XX years, range XX–XX) and XX reported gender (XX women, XX men). All formal sessions ran on the 240 Hz laboratory display with a measured refresh rate of at least 200 Hz.

**Results:** Table 17 and Fig. 9 summarize the typing task. Relative to the control condition, the readability-priority full



**FIGURE 9. Within-subject typing performance under the control and readability-priority masked conditions: each gray line is one participant (mean of two trials per condition); markers show group means with participant-level bootstrap 95% confidence intervals. Panels show transcription accuracy, words per minute, and first-keystroke latency.**



**FIGURE 10. Mean 1–5 Likert ratings (error bars: SD) for readability, stability, and immediate visual comfort across the six rated display conditions, ordered from the unmasked anchor to the readability-priority full configuration. Higher is better on every item.**

configuration changed mean accuracy by XX percentage points (95% CI [XX, XX],  $p = XX$ , effect size XX), WPM by XX (95% CI [XX, XX],  $p = XX$ , effect size XX), and first-keystroke latency by XX ms (95% CI [XX, XX],  $p = XX$ , effect size XX); CPM and attempted characters gave  $p = XX$  and XX, respectively. Table 18 and Fig. 10 summarize the subjective ratings: Friedman tests yielded  $\chi^2 = XX$  ( $p = XX$ ) for readability,  $\chi^2 = XX$  ( $p = XX$ ) for stability, and  $\chi^2 = XX$  ( $p = XX$ ) for immediate visual comfort; Holm-corrected pairwise comparisons between the readability-priority full configuration and the unmasked anchor gave  $p = XX$ , XX, and XX on the three items, respectively.

## VI. DISCUSSION

### A. PRINCIPAL FINDINGS AND RESEARCH VALUE

At the common S600 setting, readability-priority recovery is 17.8%, compared with 94.5% for unprotected content. High-suppression further lowers recovery to 5.6%. The paired content-cluster gaps remain 76.7 and 88.9 percentage points

across 288 matched units per profile. A reproducible preprocessing grid raises protected recovery, yet large paired gaps remain. Detection, tracking, VLM, integration, and simulation experiments then map how attacker capability changes these outcomes.

Relative to prior temporal-display studies, the 10,575-capture physical archive is new. So is the link from average OCR recovery to explicit-field leakage and upper-tail risk, and the evaluation of attacker escalation across preprocessing, sustained capture, and modern recognition models.

### B. INTEGRATION AND ATTACKER ESCALATION

Temporal masking exploits the difference between human temporal integration and short camera exposure. In the ideal linear model, registered complementary subframes can reconstruct the original because  $\sum \mathbf{I}_k = \mathbf{I}$ . The digital attack evaluation confirms this boundary, reaching 95.2% per-sample strongest recovery and 94.3% under pure full-cycle averaging.

Physical temporal averaging also increases recovery without requiring the mask seed or phase. The 150-frame mean reaches 71.1% for readability-priority and 47.9% for high-suppression. At the common settings, these values rise to 79.7% and 53.9%. Sustained capture thus converts temporal fragments into a usable recognition signal.

### C. PROFILE TRADE-OFFS AND PRACTICAL IMPLICATIONS

The two principal profiles occupy distinct operating points. Readability-priority preserves a digitally reconstructed SSIM of 0.948 but retains 17.7% explicit-field recovery. High-suppression lowers several recovery measures while reducing digital SSIM to 0.764 and increasing  $\Delta E_{00}$  from 1.46 to 4.87. The privacy–readability trade-off between these profiles is explicit and tunable.

The high-suppression inversion slot also reaches 93.0% recovery after contrast stretching, making slot-level composition a weak point in the current design. A practical system should therefore combine temporal masking with field-level policies and independent access controls for high-value content.

### D. VLM-AWARE IMPLICATIONS

Modern VLMs recover far more text than conventional OCR from the same protected photographs. Recovery is strongest for large-font credentials, code, digits, and short sentences; dense CET6 pages remain more strongly suppressed. The implication is that temporal profiles should adapt to glyph scale, semantic predictability, and field sensitivity.

### E. ETHICS STATEMENT

This research is defensive in nature, aiming to protect screen content from unauthorized photography. The attack analysis (§V-E2) is intended for honest assessment of security boundaries, not as an attack guide. The planned user study would involve human subjects in minimal-risk research; no participant data have been collected. Recruitment will begin only after

the competent institutional or departmental authority issues a written approval, exemption, or waiver determination. The protocol specifies informed consent, photosensitivity screening, and voluntary withdrawal as described in §V-F.

### F. LIMITATIONS AND FUTURE WORK

The physical evidence comes from one S600 UVC webcam, and profile order was fixed rather than randomized. Exposure values are control-derived labels rather than photometric shutter measurements. The archive also lacks decoded auto-exposure state, gain, white-balance, focus, ambient illuminance, display–camera phase, and physical panel luminance. These factors can interact with time, geometry, and profile order.

The current PoC operates at the application layer with fixed panel output. It lacks photometric slot-timing measurements, luminance-matched static controls, and matched noise-off versions of the enhanced profiles. Temporal sparsity, duty-cycle luminance, overlays, panel response, camera exposure, and ISP processing therefore remain coupled. Rolling-shutter row mixing and panel transitions are candidate explanations for the simulation-to-real gap.

The content-cluster intervals hold the tested geometries fixed, and the 12-item corpus is too small for script-level inference. Detection uses 150 COCO images at one position, while tracking uses 450 still-frame captures from one MOT17 sequence. The VLM probes use one on-axis profile, three API endpoints, and raw JPEG inputs. Commercial OCR, learned restoration, and unconstrained adaptive preprocessing may recover additional text. The 150-frame experiment also leaves the duration–recovery curve unresolved.

The planned user study has not yet collected participant data. It evaluates readability-priority rather than high-suppression, and its immediate-comfort rating does not address long-term or clinical effects. Driver-level deployment would also require framebuffer integration, panel coordination, and verified timing. Temporal masking should therefore complement established access controls when protecting credentials and other high-value fields.

Future work should randomize profile order, lock camera parameters, record batch time, and measure physical slot timing and luminance. Matched static controls can separate temporal and luminance effects. Replication across cameras and displays can test transfer, while continuous-video attacks can map recovery against recording duration. VLM-aware profiles can explore font-size-adaptive granularity, glyph decoys, and field-level rendering. Completing the preregistered usability study would connect machine-recovery gains to human readability and comfort.

## VII. CONCLUSION

We evaluated temporal pixel masking on one UVC camera link. Across 8 common-setting geometries, the same 288 matched units per profile (12 content items  $\times$  8 geometries  $\times$  3 repeats) yield 94.5% mean recovery unprotected, 17.8% for readability-priority, and 5.6% for high-suppression. Paired

content-cluster contrasts are 76.7 percentage points ([74.3, 79.2]) and 88.9 points ([85.5, 91.9]). Readability-priority retains 17.7% matched field-micro recovery and 95.0% P99 character recovery, while high-suppression misses its descriptive  $< 5\%$  target.

Attack-oriented evaluation recovers substantially more information. A 150-frame temporal mean reaches 71.1% character recovery for readability-priority and 47.9% for high-suppression; the common-setting values are 79.7% and 53.9%. In the full 9-geometry sensitivity pool, long-exposure recovery is 60.9% for readability-priority versus 47.3% unprotected. On-axis VLM probes also recover substantial large-font content, while the three-engine preprocessing oracle reaches 40.2% for readability-priority and 13.7% for high-suppression.

Detection and tracking stress tests show the same directional change across four detectors. The archived data, scripts, and per-sample records are public and can serve as a baseline for comparing future temporal display-privacy methods and attacks.

## ACKNOWLEDGMENT

### [PLACEHOLDER: Acknowledgment to be filled.]

OpenAI Codex [56] assisted with analysis-script development, consistency checks, and English revision of the Abstract. The authors specified the scientific decisions, verified generated analyses against archived data, reviewed all retained text and code, and remain responsible for the manuscript.

## REFERENCES

- [1] X. Wang, K. Peng, X. Yi, and H. Li, "Mind the gap: Mapping wearer-bystander privacy tensions and context-adaptive pathways for camera glasses," in *Proc. CHI Conf. Human Factors in Computing Systems*, 2026, pp. 1–28, doi: 10.1145/3772318.3791848. [Online]. Available: <https://doi.org/10.1145/3772318.3791848>
- [2] Ponemon Institute, "Global visual hacking experiment," 3M, Tech. Rep., 2016, Accessed: Jul. 2, 2026. [Online]. Available: [https://www.3m.com/3M/en\\_US/privacy-screen-protectors-us/expertise/visualhacking/](https://www.3m.com/3M/en_US/privacy-screen-protectors-us/expertise/visualhacking/)
- [3] M. Eiband, M. Khamis, E. von Zezschwitz, H. Hussmann, and F. Alt, "Understanding shoulder surfing in the wild: Stories from users and observers," in *Proc. CHI Conf. Human Factors in Computing Systems*, 2017, pp. 4254–4265, doi: 10.1145/3025453.3025636.
- [4] M. Backes, T. Chen, M. Dürmuth, H. P. A. Lensch, and M. Welk, "Tempest in a teapot: Compromising reflections revisited," in *Proc. IEEE Symp. Security and Privacy (S&P)*, 2009, pp. 315–327, doi: 10.1109/SP.2009.20.
- [5] S. Fernández, E. Martínez, G. Varela, P. Musé, and F. Larroca, "Deep-TEMPEST: Using deep learning to eavesdrop on HDMI from its unintended electromagnetic emanations," in *Proc. 13th Latin-American Symp. Dependable and Secure Computing (LADC)*, 2024, pp. 91–100, doi: 10.1145/3697090.3697094.
- [6] R. Raguram, A. M. White, D. Goswami, F. Monrose, and J.-M. Frahm, "iSpy: Automatic reconstruction of typed input from compromising reflections," in *Proc. 18th ACM Conf. Computer and Communications Security (CCS)*, 2011, pp. 527–536, doi: 10.1145/2046707.2046769.
- [7] B. Tang and K. G. Shin, "Eye-Shield: Real-time protection of mobile device screen information from shoulder surfing," in *Proc. 32nd USENIX Security Symp.*, 2023, pp. 6695–6712.
- [8] C.-Y. Chen, B.-Y. Lin, J. Wang, and K. G. Shin, "Keep Others from Peeking at Your Mobile Device Screen!" in *Proc. 25th Annu. Int. Conf. Mobile Computing and Networking (MobiCom)*, 2019, pp. 1–16, doi: 10.1145/3300061.3300119.
- [9] S. Zhu, C. Zhang, and X. Zhang, "Automating visual privacy protection using a smart LED," in *Proc. 23rd ACM Int. Conf. Mobile Computing and Networking (MobiCom)*, 2017, pp. 329–342, doi: 10.1145/3117811.3117820.
- [10] M. Naor and A. Shamir, "Visual cryptography," in *Proc. EUROCRYPT*, 1994, pp. 1–12.
- [11] X. Zhong, A. Das, F. Alrasheedi, and A. Tanvir, "A brief, in-depth survey of deep learning-based image watermarking," *Applied Sciences*, vol. 13, no. 21, p. 11852, 2023, doi: 10.3390/app132111852.
- [12] H. Fang, W. Zhang, H. Zhou, H. Cui, and N. Yu, "Screen-shooting resilient watermarking," *IEEE Trans. Information Forensics and Security*, vol. 14, no. 6, pp. 1403–1418, 2019.
- [13] V. Nguyen, Y. Tang, A. Ashok, M. Gruteser, K. Dana, W. Hu, E. Wengrowski, and N. Mandayam, "High-rate flicker-free screen-camera communication with spatially adaptive embedding," in *Proc. IEEE INFOCOM*, 2016, pp. 1–9, doi: 10.1109/INFOCOM.2016.7524512.
- [14] V. Tran, G. Jayatilaka, A. Ashok, and A. Misra, "DeepLight: Robust & unobtrusive real-time screen-camera communication for real-world displays," in *Proc. 20th ACM/IEEE Int. Conf. Information Processing in Sensor Networks (IPSN)*, 2021, pp. 238–253, doi: 10.1145/3412382.3458269.
- [15] S. K. Ghiasi, M. Kaldenbach, and M. Zuniga, "Passive screen-to-camera communication," in *Proc. 20th Int. Conf. Distributed Computing in Smart Systems and the Internet of Things (DCOSS-IoT)*, 2024, pp. 35–43, doi: 10.1109/DCOSS-IoT61029.2024.00016.
- [16] W. S. Yerazunis and M. Carbone, "Privacy-enhanced displays by time-masking images," in *Proc. Australian Conf. Computer-Human Interaction (OzCHI)*, Nov. 2001. [Online]. Available: <https://www.merl.com/publications/TR2002-11>
- [17] X. Wu and G. Zhai, "Temporal psychovisual modulation: A new paradigm of information display," *IEEE Signal Process. Mag.*, vol. 30, no. 1, pp. 136–141, 2013, doi: 10.1109/MSP.2012.2219678.
- [18] Z. Gao, G. Zhai, and X. Min, "Information security display system based on temporal psychovisual modulation," in *Proc. IEEE Int. Symp. Circuits and Systems (ISCAS)*, 2014, pp. 449–452, doi: 10.1109/IS-CAS.2014.6865167.
- [19] L. Zhang, C. Bo, J. Hou, X.-Y. Li, Y. Wang, K. Liu, and Y. Liu, "Kaleido: You can watch it but cannot record it," in *Proc. 21st Annu. Int. Conf. Mobile Computing and Networking (MobiCom)*, 2015, pp. 372–385, doi: 10.1145/2789168.2790106.
- [20] L. Yang, W. Wang, Z. Wang, and Q. Zhang, "Rainbow: Preventing mobile-camera-based piracy in the physical world," in *Proc. IEEE INFOCOM*, 2018, pp. 1061–1069, doi: 10.1109/INFOCOM.2018.8485881.
- [21] J. Lim, "Defeat spyware with anti-screen capture technology using visual persistence," in *Proc. 3rd Symp. Usable Privacy and Security (SOUPS)*, 2007, pp. 147–148, doi: 10.1145/1280680.1280701.
- [22] I. J. Goodfellow, J. Shlens, and C. Szegedy, "Explaining and harnessing adversarial examples," in *Proc. Int. Conf. Learning Representations (ICLR)*, 2015.
- [23] A. Madry, A. Makelov, L. Schmidt, D. Tsipras, and A. Vladu, "Towards deep learning models resistant to adversarial attacks," in *Proc. Int. Conf. Learning Representations (ICLR)*, 2018.
- [24] N. Carlini and D. Wagner, "Towards evaluating the robustness of neural networks," in *Proc. IEEE Symp. Security and Privacy (S&P)*, 2017, pp. 39–57.
- [25] K. Eykholt, I. Evtimov, E. Fernandes, B. Li, A. Rahmati, C. Xiao, A. Prakash, T. Kohno, and D. Song, "Robust physical-world attacks on deep learning visual classification," in *Proc. IEEE/CVF Conf. Computer Vision and Pattern Recognition (CVPR)*, 2018, pp. 1625–1634.
- [26] S. Thys, W. Van Ranst, and T. Goedeme, "Fooling automated surveillance cameras: Adversarial patches to attack person detection," in *Proc. IEEE/CVF Conf. Computer Vision and Pattern Recognition Workshops (CVPRW)*, 2019.
- [27] K. Xu, G. Zhang, S. Liu, Q. Fan, M. Sun, H. Chen, P.-Y. Chen, Y. Wang, and X. Lin, "Adversarial T-shirt! Evading person detectors in a physical world," in *Proc. European Conf. Computer Vision (ECCV)*, 2020, pp. 665–681.
- [28] H. Wei, H. Tang, X. Jia, Z. Wang, H. Yu, Z. Li, S. Satoh, L. Van Gool, and Z. Wang, "Physical adversarial attack meets computer vision: A decade survey," *IEEE Trans. Pattern Anal. Mach. Intell.*, vol. 46, no. 12, pp. 9797–9817, 2024.
- [29] A. Athalye, L. Engstrom, A. Ilyas, and K. Kwok, "Synthesizing robust adversarial examples," in *Proc. Int. Conf. Machine Learning (ICML)*, 2018, pp. 284–293.

- [30] J. Gu, X. Jia, P. de Jorge, W. Yu, X. Liu, A. Ma, Y. Xun, A. Hu, A. Khakzar, Z. Li, X. Cao, and P. Torr, "A survey on transferability of adversarial examples across deep neural networks," *Transactions on Machine Learning Research*, 2024, [Online]. Available: <https://openreview.net/forum?id=AYJ3m7BocI>.
- [31] D. Niu, R. Guo, and Y. Wang, "Moiré attack (MA): A new potential risk of screen photos," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 34, 2021, pp. 25 543–25 554.
- [32] D. J. Bernstein, "ChaCha, a variant of Salsa20," in *Workshop Record of SASC*, vol. 8, 2008, pp. 3–5.
- [33] D. E. Knuth, *The Art of Computer Programming, Volume 2: Seminumerical Algorithms*, 3rd ed. Addison-Wesley, 1997.
- [34] C. Song and V. Shmatikov, "Fooling OCR systems with adversarial text images," *arXiv preprint arXiv:1802.05385*, 2018, doi: 10.48550/arXiv.1802.05385. [Online]. Available: <https://arxiv.org/abs/1802.05385>
- [35] G. Sharma, W. Wu, and E. N. Dalal, "The CIEDE2000 color-difference formula: Implementation notes, supplementary test data, and mathematical observations," *Color Research & Application*, vol. 30, no. 1, pp. 21–30, 2005.
- [36] Z. Wang, A. C. Bovik, H. R. Sheikh, and E. P. Simoncelli, "Image quality assessment: From error visibility to structural similarity," *IEEE Trans. Image Process.*, vol. 13, no. 4, pp. 600–612, 2004.
- [37] Y. Cai, A. Bozorgian, M. Ashraf, R. Wanat, and R. K. Mantiuk, "elaTCSF: A temporal contrast sensitivity function for flicker detection and modeling variable refresh rate flicker," in *Proc. SIGGRAPH Asia*, 2024.
- [38] J. Davis, Y.-H. Hsieh, and H.-C. Lee, "Humans perceive flicker artifacts at 500 Hz," *Scientific Reports*, vol. 5, 2015, Art. no. 7861, doi: 10.1038/srep07861.
- [39] EMEET, "SmartCam S600: 4K Ultra HD webcam," 2026, Accessed: Jul. 7, 2026. [Online]. Available: <https://emeet.com/en-eu/products/webcam-s600>
- [40] V. Paruchuri, "Surya: Multilingual document OCR toolkit," 2024. [Online]. Available: <https://github.com/VikParuchuri/surya>
- [41] JaidevAI, "EasyOCR: Ready-to-use OCR with 80+ supported languages," 2020. [Online]. Available: <https://github.com/JaidevAI/EasyOCR>
- [42] R. Smith, "An overview of the Tesseract OCR engine," in *Proc. Int. Conf. Document Analysis and Recognition (ICDAR)*, 2007, pp. 629–633.
- [43] S. Bai et al., "Qwen3-VL Technical Report," *arXiv preprint arXiv:2511.21631*, 2025, doi: 10.48550/arXiv.2511.21631. [Online]. Available: <https://arxiv.org/abs/2511.21631>
- [44] Moonshot AI, "Kimi-K2.6 model card," Hugging Face, 2026, accessed: Jul. 11, 2026. [Online]. Available: <https://huggingface.co/moonshotai/Kimi-K2.6>
- [45] GLM-V Team, W. Hong, W. Yu, X. Gu et al., "GLM-4.1V-Thinking and GLM-4.5V: Towards Versatile Multimodal Reasoning with Scalable Reinforcement Learning," *arXiv preprint arXiv:2507.01006*, 2025, doi: 10.48550/arXiv.2507.01006. [Online]. Available: <https://arxiv.org/abs/2507.01006>
- [46] Y. Shi, D. Peng, W. Liao, Z. Lin, X. Chen, C. Liu, Y. Zhang, and L. Jin, "Exploring OCR capabilities of GPT-4V(ision): A quantitative and in-depth evaluation," *arXiv preprint arXiv:2310.16809*, 2023, doi: 10.48550/arXiv.2310.16809. [Online]. Available: <https://arxiv.org/abs/2310.16809>
- [47] C. Jiang, Z. Wang, M. Dong, and J. Gui, "Survey of adversarial robustness in multimodal large language models," *arXiv preprint arXiv:2503.13962*, 2025, doi: 10.48550/arXiv.2503.13962. [Online]. Available: <https://arxiv.org/abs/2503.13962>
- [48] G. Jocher, J. Qiu, M. Liu, S. Lyu, F. C. Akyon, and M. E. Kalfaoglu, "Ultralytics YOLO26: Unified real-time end-to-end vision models," 2026, arXiv:2606.03748, doi: 10.48550/arXiv.2606.03748. [Online]. Available: <https://arxiv.org/abs/2606.03748>
- [49] Y. Zhao, W. Lv, S. Xu, J. Wei, G. Wang, Q. Dang, Y. Liu, and J. Chen, "DETRs beat YOLOs on real-time object detection," in *Proc. IEEE/CVF Conf. Computer Vision and Pattern Recognition (CVPR)*, 2024, pp. 16 965–16 974.
- [50] S. Ren, K. He, R. Girshick, and J. Sun, "Faster R-CNN: Towards real-time object detection with region proposal networks," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 28, 2015, pp. 91–99.
- [51] T.-Y. Lin, P. Goyal, R. Girshick, K. He, and P. Dollár, "Focal loss for dense object detection," in *Proc. IEEE Int. Conf. Computer Vision (ICCV)*, 2017, pp. 2980–2988.
- [52] Y. Zhang, P. Sun, Y. Jiang, D. Yu, F. Weng, Z. Yuan, P. Luo, W. Liu, and X. Wang, "ByteTrack: Multi-object tracking by associating every detection box," in *Proc. European Conf. Computer Vision (ECCV)*, 2022, pp. 1–21.
- [53] T.-Y. Lin, M. Maire, S. Belongie, J. Hays, P. Perona, D. Ramanan, P. Dollár, and C. L. Zitnick, "Microsoft COCO: Common objects in context," in *Proc. European Conf. Computer Vision (ECCV)*, 2014, pp. 740–755.
- [54] A. Milan, L. Leal-Taixé, I. Reid, S. Roth, and K. Schindler, "MOT16: A benchmark for multi-object tracking," *arXiv preprint arXiv:1603.00831*, 2016.
- [55] R. W. Soukoreff and I. S. MacKenzie, "Metrics for text entry research: An evaluation of MSD and KSPC, and a new unified error metric," in *Proc. SIGCHI Conf. Human Factors in Computing Systems (CHI)*, 2003, pp. 113–120, doi: 10.1145/642611.642632.
- [56] OpenAI, "Introducing the Codex App," OpenAI, 2026, accessed: Jul. 10, 2026. [Online]. Available: <https://openai.com/index/introducing-the-codex-app/>

...